

Convolutional Neural Network

Reference: <https://www.coursera.org/learn/convolutional-neural-networks/home/welcome>

© Ronak Patel, Computer Engineering Department , CSPIT, CHARUSAT

Application



64x64

→ Cat? (0/1)

Image Classification



Object Detection



Neural Style transfer

Why not NN?



$64 \times 64 \times 3$


→ Cat? (0/1)

12288



$1000 \times 1000 \times 3$
= 3 million


Convolution



6×6

$n \times n$


*



3×3

$f \times f$

=




4×4

$(n - f + 1) \times (n - f + 1)$

Vertical edge detection

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	<u>10</u>	<u>10</u>	<u>0</u>	0	0
10	<u>10</u>	<u>10</u>	<u>0</u>	0	0
10	<u>10</u>	<u>10</u>	<u>0</u>	0	0

6x6




*

1	0	-1
1	0	-1
1	0	-1

3x3


*



=

0	30	30	0
0	30	30	0
0	30	30	0
0	<u>30</u>	<u>30</u>	0

4x4




Padding

Valid Conv. = no padding

Same Conv. = pad so that output size is the same as input size.

$$p = (f - 1) / 2$$




$$n \times n$$

$$8 \times 8$$

Padding (p) = 1


*



$$f \times f$$

$$3 \times 3$$

=




$$6 \times 6$$

$$(n + 2p - f + 1) \times (n + 2p - f + 1)$$

Why padding?

- Throwing away information from edges.
- Shrink the image every time after convolution.


Strided Convolution



$n \times n$
 8×8


Padding (p) = 1
Stride (S) = 2

*



$f \times f$
 3×3


=




6×6

$$\left\lfloor \frac{n + 2p - f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n + 2p - f}{s} + 1 \right\rfloor$$

Convolution on RGB image






Summary: $n \times n \times n_c \times f \times f \times n_c \rightarrow \frac{n-f+1}{4} \times \frac{n-f+1}{4} \times n'_c$

$6 \times 6 \times 3 \times 3 \times 3 \times 3 \rightarrow 4 \times 4 \times 2 \uparrow \# \text{ filters}$

Andrew Ng




$6 \times 6 \times 3$

$a^{[0]}$


$$z^{[1]} = w^{[1]} a^{[0]} + b^{[1]}$$

$$a^{[1]} = g(z^{[1]})$$

$*$




$3 \times 3 \times 3$




$3 \times 3 \times 3$

$w^{[1]}$



$6 \times 6 \times 3$

$a^{[0]}$



Summary of Notation

If layer l is a convolution layer:

$f^{[l]}$ = filter size

$p^{[l]}$ = padding

$s^{[l]}$ = stride

$n_c^{[l]}$ = number of filters

→ Each filter is: $f^{[l]} \times f^{[l]} \times n_c^{[l-1]}$

Activations: $a^{[l]} \rightarrow n_H^{[l]} \times n_W^{[l]} \times n_c^{[l]}$

Weights: $f^{[l]} \times f^{[l]} \times n_c^{[l-1]} \times n_c^{[l]}$

bias: $n_c^{[l]} - (1, 1, 1, n_c^{[l]})$ ← #f: filters in layer l.

Input: $n_H^{[l-1]} \times n_W^{[l-1]} \times n_c^{[l-1]}$ ←
Output: $n_H^{[l]} \times n_W^{[l]} \times n_c^{[l]}$ ←


$$n_{HW}^{[l]} = \left\lfloor \frac{n_H^{[l-1]} + 2p^{[l]} - f^{[l]}}{s^{[l]}} + 1 \right\rfloor$$

$$A^{[l]} \rightarrow m \times \underbrace{n_H^{[l]} \times n_W^{[l]} \times n_c^{[l]}}_{n_c^{[l]} \times n_H^{[l]} \times n_W^{[l]}}$$

ConvNet

	Shape
Input	(39,39,3)
Conv1 (f=3, s=1, p=0), 10-filters	
Conv2 (f=5, s=2, p=0) 20-filters	
Conv3 (f=5, s=2, p=0) 40-filters	
Flatten	

Example ConvNet




Andrew Ng


Types of Layer in ConvNets

- Convolution (Conv)
- Pooling (pool)
- Fully Connected (FC)

Pooling

Max pooling





No parameters to learn in pooling

Summary of pooling

Hyperparameters:

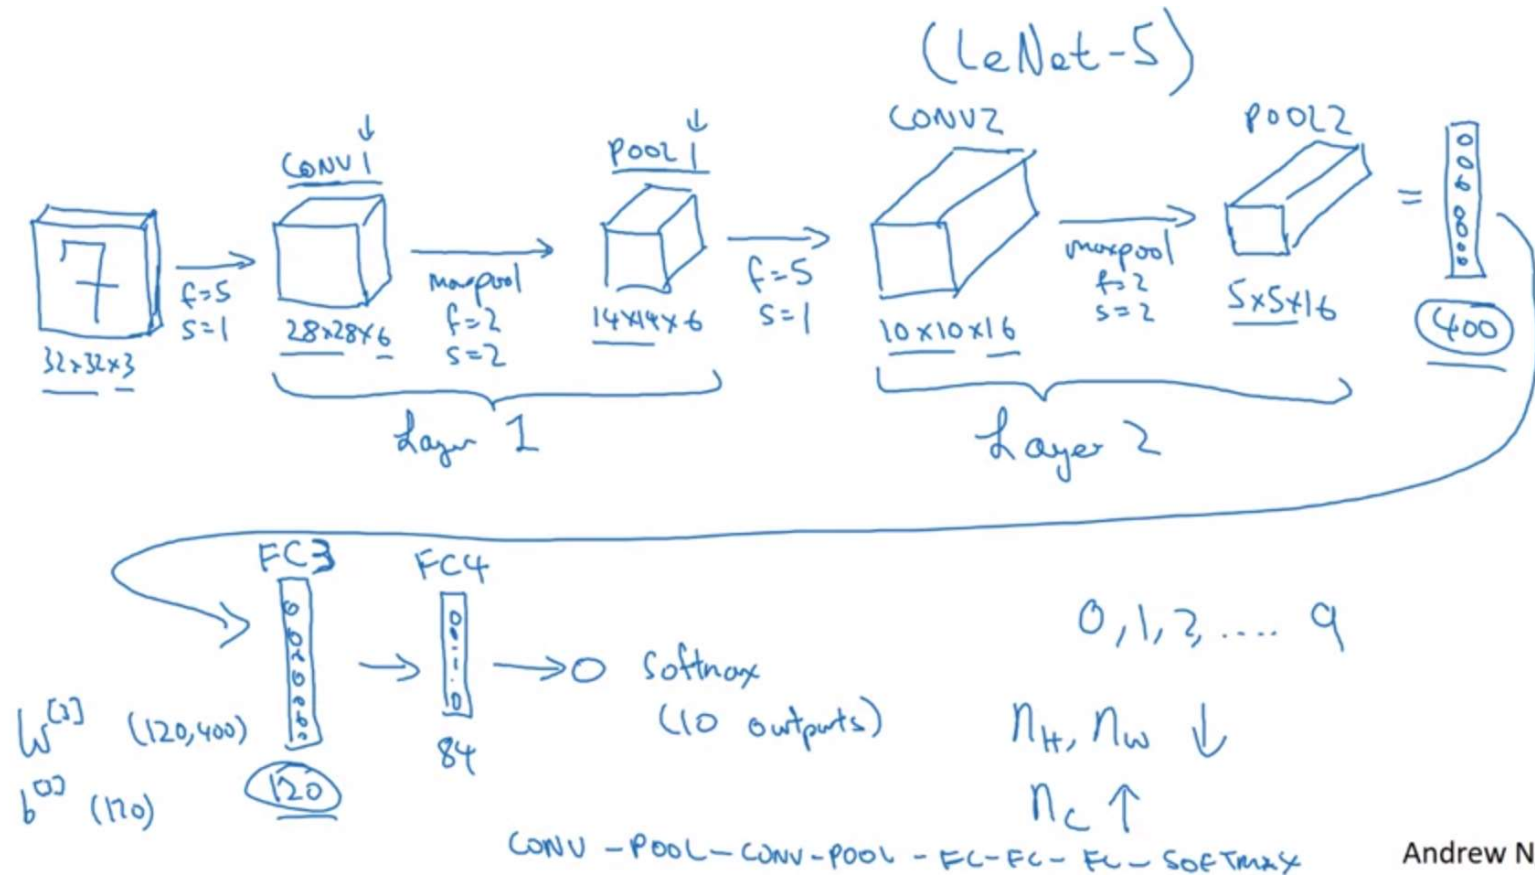
f : filter size $f=2, s=2$
 s : stride $f=3, s=2$
Max or average pooling

~~$\Rightarrow p$: padding.~~

No parameters to learn!

$$\begin{array}{c} n_H \times n_W \times \underline{n_C} \\ \downarrow \\ \left\lfloor \frac{n_H - f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n_W - f}{s} + 1 \right\rfloor \\ \times \underline{n_C} \end{array}$$

LeNet-5



	Shape	Parameters
Input	(32,32,3)	-
Conv1 (f=5, s=1), 8-filters (3D-Filters)	(28,28,8)	608
Pool1	(14,14,8)	-
Conv2 (f=5, s=1) 16-filters	(10,10,16)	3216
Pool2	(5,5,16)	-
FC3	120	48120
FC4	84	10164
Output-SoftMax	10	850

List of Parameters and Hyperparameters

Hyperparameters: Which needs to be tuned, Usually with validation set

1. Kernel size - 3 x 3, 5 x 5, 7 x 7
2. Stride - 1, 2 or 3
3. Padding size - 0, 1, 2
4. Number of filters - 32, 64, 128, etc..
5. No of hidden layers - Trial and error
6. No of neurons in each hidden layer - Trial and error
7. Learning rate- 0.1, 0.01, 0.001
8. Activation functions
9. Number of iteration – Trial and error
10. No. of Epoch – Trial and error
11. Batch size – 4, 8, 16, 32, 64, 128 (Power of 2)
12. Regularization – Dropout, L1, L2
13. Weights and Bias Initialization – Zero, Random, He
14. Optimization algorithm – SGD, ADAM, RMSProp
15. Learning Rate decay
16. Momentum – 0.9
17. Number of iterations, epoch and batch size

Parameters: Which are learned during training, Weights and bias.

Transfer Learning

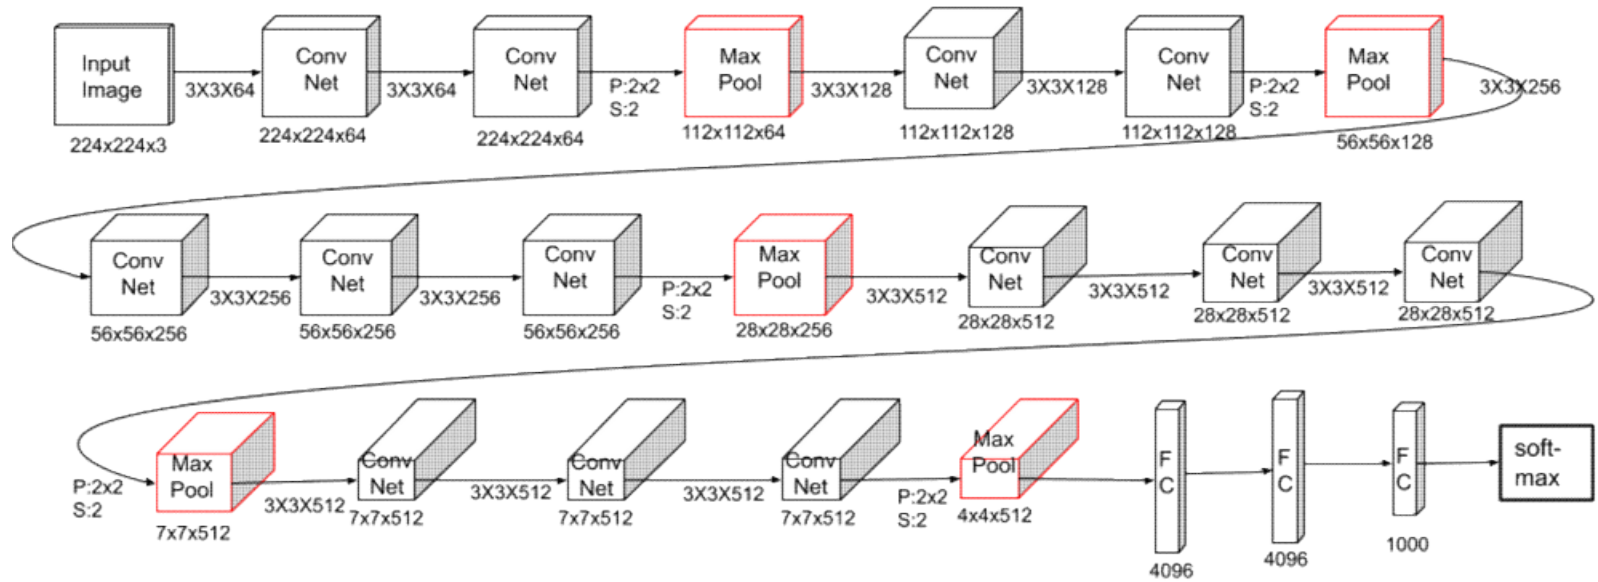
- Transfer learning is a machine learning technique where a model trained on one task is re-purposed on a second related task.
- Transfer learning with convolutional neural networks (CNNs) has revolutionized the field of computer vision by enabling the reuse of pre-trained models on new, related tasks. This powerful technique leverages the knowledge learned from large-scale datasets, allowing for faster and more accurate model training, even with limited labeled data.
- transfer learning significantly reduces the need for extensive training time and computational resources.

- Transfer learning takes advantage of the fact that CNNs trained on large datasets, such as ImageNet, have learned general features that are relevant to many visual tasks. Instead of training a CNN from scratch on a new dataset, transfer learning involves using a pre-trained CNN as a starting point and fine-tuning it on the new dataset.
- The pre-trained CNN acts as a feature extractor, capturing high-level visual representations. These features are then passed to new layers designed for the specific task. The pre-trained layers are frozen during fine-tuning, while the new layers are adjusted.

Steps to Implement Transfer Learning

- Select a Pre-trained Model
- Load Pre-trained Model
- Customize the Model
- Freeze Pre-trained Layers
- Prepare Data
- Train the Model
- Fine-tuning (Optional)
- Evaluate and Test

VGG-16



VGG-16



Example for transfer learning

```
import tensorflow as tf
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Flatten, Dense
from tensorflow.keras.utils import to_categorical

# Load CIFAR-10 dataset
(itrain, ltrain), (itest, ltest) = cifar10.load_data()

# Preprocess the data
itrain = itrain / 255.0
itest = itest / 255.0
ltrain = to_categorical(ltrain)
ltest = to_categorical(ltest)
```

```

# Load pre-trained VGG16 model (excluding the top fully-connected layers)
basem = VGG16(weights='imagenet', include_top=False, input_shape=(32, 32, 3))

# Freeze the pre-trained layers
for layer in basem.layers:
    layer.trainable = False

# Create a new model on top
semodel = Sequential()
semodel.add(basem)
semodel.add(Flatten())
semodel.add(Dense(256, activation='relu'))
semodel.add(Dense(10, activation='softmax')) # CIFAR-10 has 10 classes

# Compile the model
semodel.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Train the model
semodel.fit(itrain, ltrain, epochs=10, batch_size=32, validation_data=(itest, ltest))

# Evaluate the model on test data
ltest, atest = semodel.evaluate(itest, ltest)
print("Test accuracy:", atest)

```



Any
Questions

